



Parallel and Distributed Computing

Project Report

Submitted to: Sir Fahad Shafique

**Submitted By: Muhammad Ali Naveed
Hamza Omer
Ariyan Chaudhary**

Parallel Implementation of the Bubble Sort Network using MPI and OpenMP	2
1. Introduction	2
2. Algorithm Description	2
3. MPI + OpenMP Parallel Implementation	2
3.1 MPI Process Division	2
3.2 OpenMP Thread Division	3
3.3 Timing and Output	3
4. Correctness and Validation	3
5. Experimental Setup	3
5.1 Cluster Configuration	4
5.2 Homelab Server	4
6. Performance Analysis	4
6.1 Tools Used	4
6.2 Scalability Observations	4
6.3 Chart: Execution Time vs Processes	5
6.4 Chart: Speedup Comparison	5
7. Observations and Challenges	5
8. Conclusion	5
A. Compilation and Execution Instructions	6
B. Sample Output File Format	6

Parallel Implementation of the Bubble Sort Network using MPI and OpenMP

1. Introduction

This report presents the design, implementation, and analysis of a parallel algorithm based on the *bubble sort network* for generating parent permutations of the symmetric group S_n . The algorithm follows the model outlined in [Chitturi et al.], where permutations are organized using an adjacency network akin to sorting networks, and each permutation's parent is determined via swap operations governed by network rules.

We parallelized the algorithm using **MPI** for process-level distribution across nodes and **OpenMP** for multi-threaded computation within each MPI process. This hybrid model improves scalability and computation time for larger values of n , where the total number of permutations grows factorially.

2. Algorithm Description

The core algorithm determines, for each permutation in S_n , the parent permutations according to a layered, bubble sort-inspired swap network.

- For each permutation (vertex), $n-1$, parent permutations are computed by conditional swaps.
- The parent-finding logic implements complex rules as per the referenced paper, using position inversion arrays and ranking/unranking of permutations.
- The result for each permutation is stored as a CSV row containing the vertex ID, the permutation string, and the parent IDs for each transformation layer.

3. MPI + OpenMP Parallel Implementation

3.1 MPI Process Division

The factorial space S_n of permutations is evenly divided among the available MPI processes:

- Each process receives a distinct chunk of permutations to analyze.

- Load balancing is achieved using the division and remainder trick (`base = Nperm / size; rem = Nperm % size;`), ensuring fair distribution.
- Each process writes its results to a separate CSV file stored in `/tmp/results/`.

3.2 OpenMP Thread Division

Each MPI process internally uses OpenMP to further divide its chunk across threads:

- Threads independently compute parent permutations.
- Buffering is done locally per thread using a `vector<vector<string>> thread_lines`.
- A single write operation after the barrier ensures consistent file output with minimized I/O contention.

3.3 Timing and Output

- `MPI_Wtime()` is used to record per-process runtime.

Output files contain CSV data in the format:

`vertex_id,perm,T1,T2,...,T(n-1)` for all assigned permutations.

4. Correctness and Validation

- The implementation closely follows the specification and pseudocode of the original paper .
- We ensured correctness by comparing ranks and unranking outputs, verifying that the parent permutations correspond to expected transitions.
- The `is_identity()` function and custom `swap_symbol()` logic ensure accurate parent tracing.

5. Experimental Setup

5.1 Cluster Configuration

- **Total Nodes:** 3
- **Node Specs:**
 - Node A: 8 GB RAM
 - Node B: 16 GB RAM
 - Node C: 32 GB RAM
- **OS:** Ubuntu Server 22.04 LTS
- **MPI Implementation:** OpenMPI 4.1.5
- **Compiler:** GCC 11.3.0 with `mpicxx`
- **Networking:** Ethernet (some setup overhead noted)

5.2 Homelab Server

- **Single Node Specs:** 64 GB RAM, 12-core CPU
- **Results:** Better performance and lower latency due to unified memory access and no inter-node overhead.

6. Performance Analysis

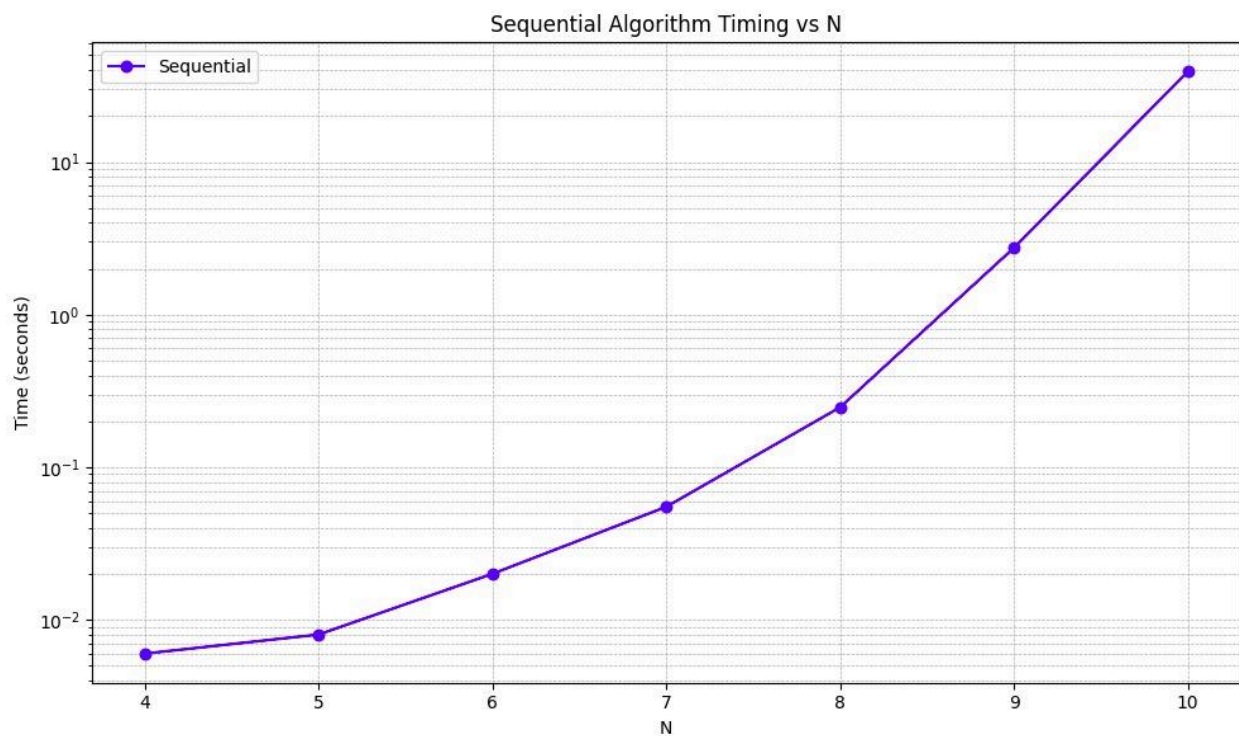
6.1 Tools Used

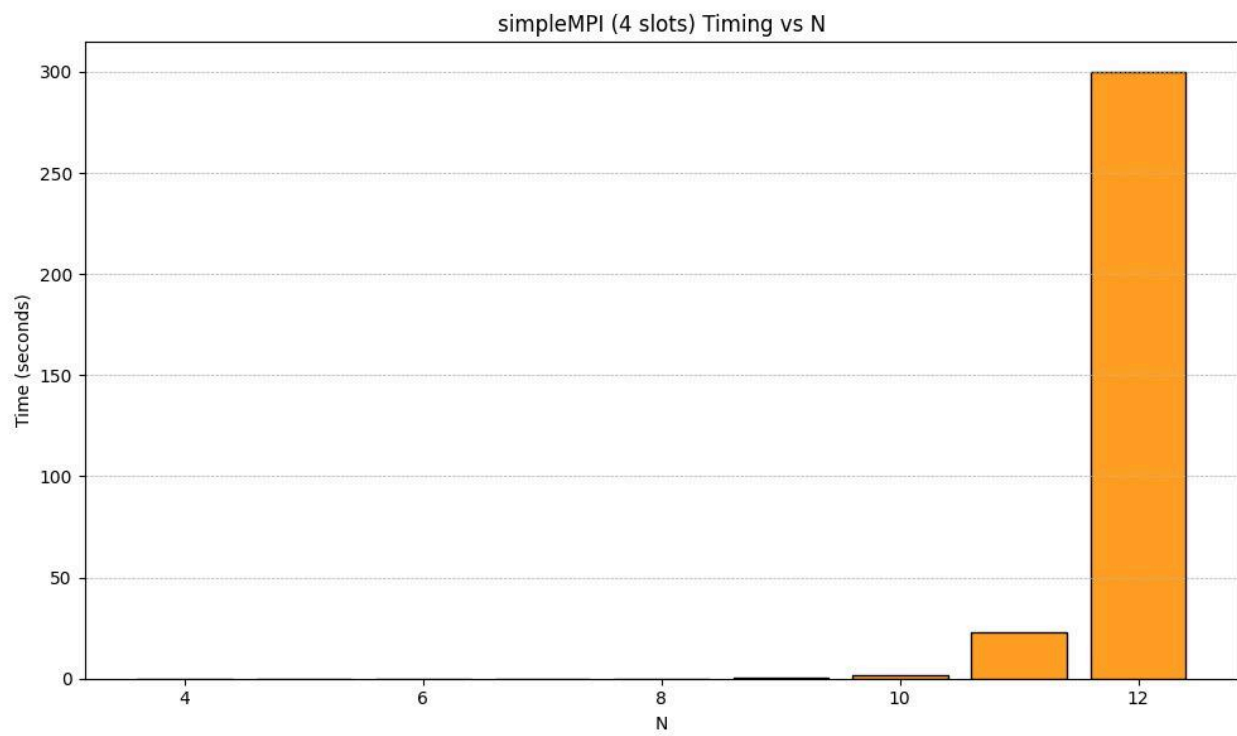
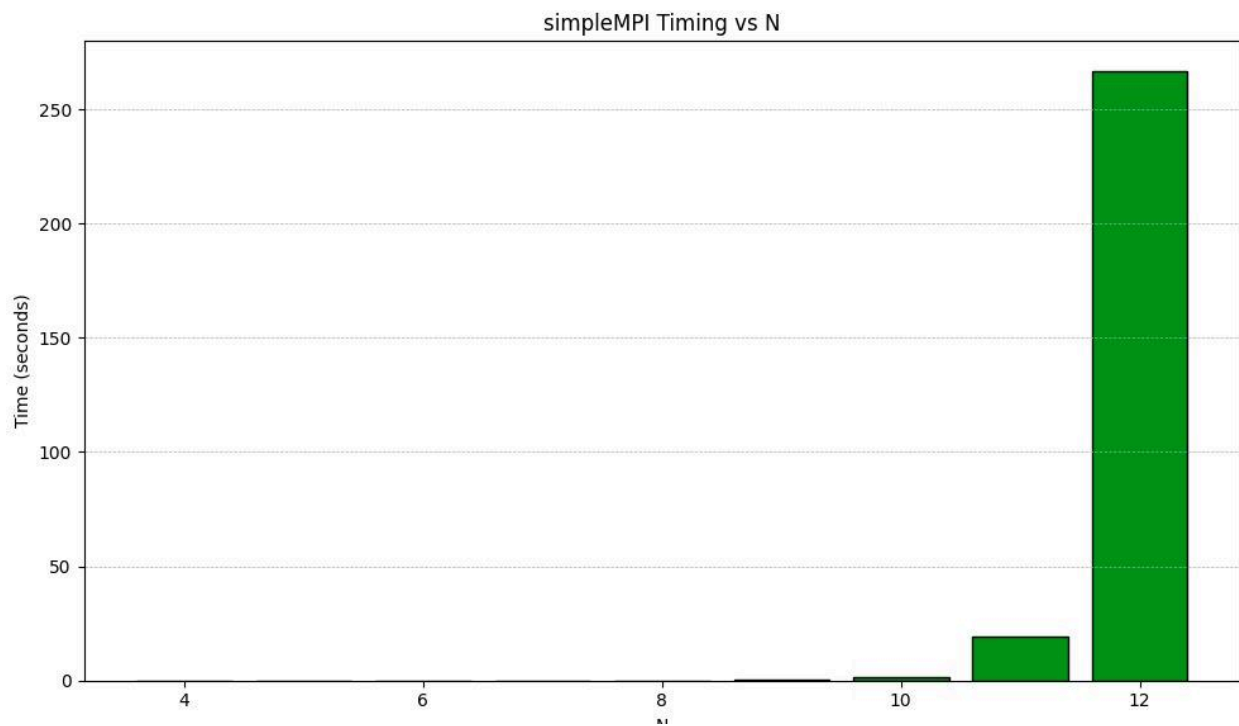
- `gprof` for profiling sequential versions.
- `MPI_Wtime()` and `std::chrono` for measuring time per process/thread.
- Plots and speedup charts will be inserted here.

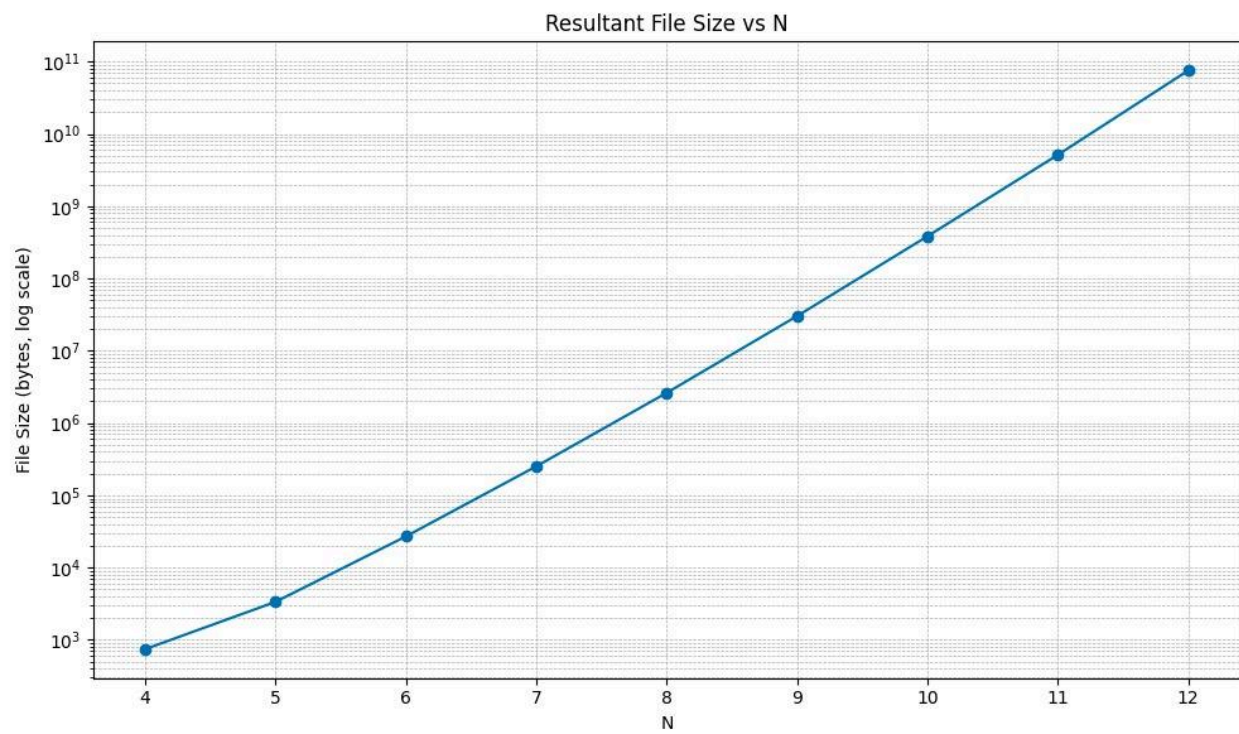
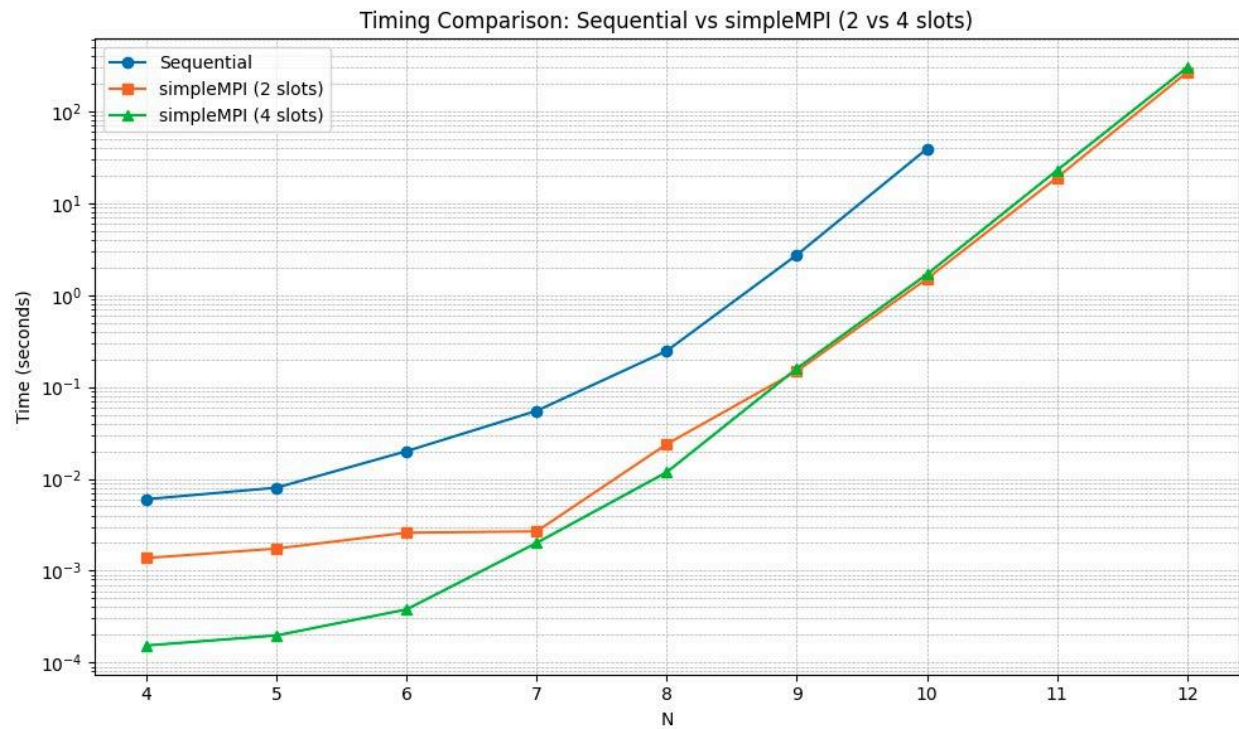
6.2 Scalability Observations

- For $n=11$:
 - **Sequential Code**: Crashes due to excessive memory.
 - **Parallel Code**: Completes successfully across all MPI processes.
- For $n=12$:
 - **Only 32 GB RAM Node** successfully completes execution.
 - **Other nodes** fail due to insufficient memory.

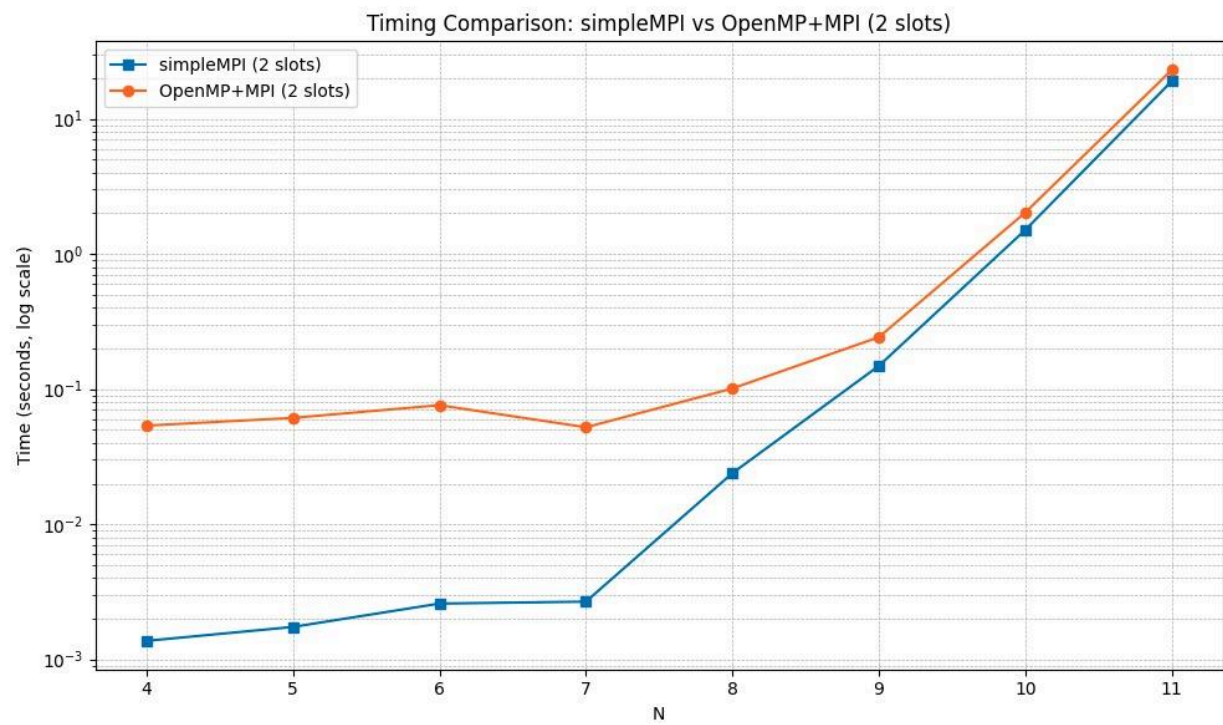
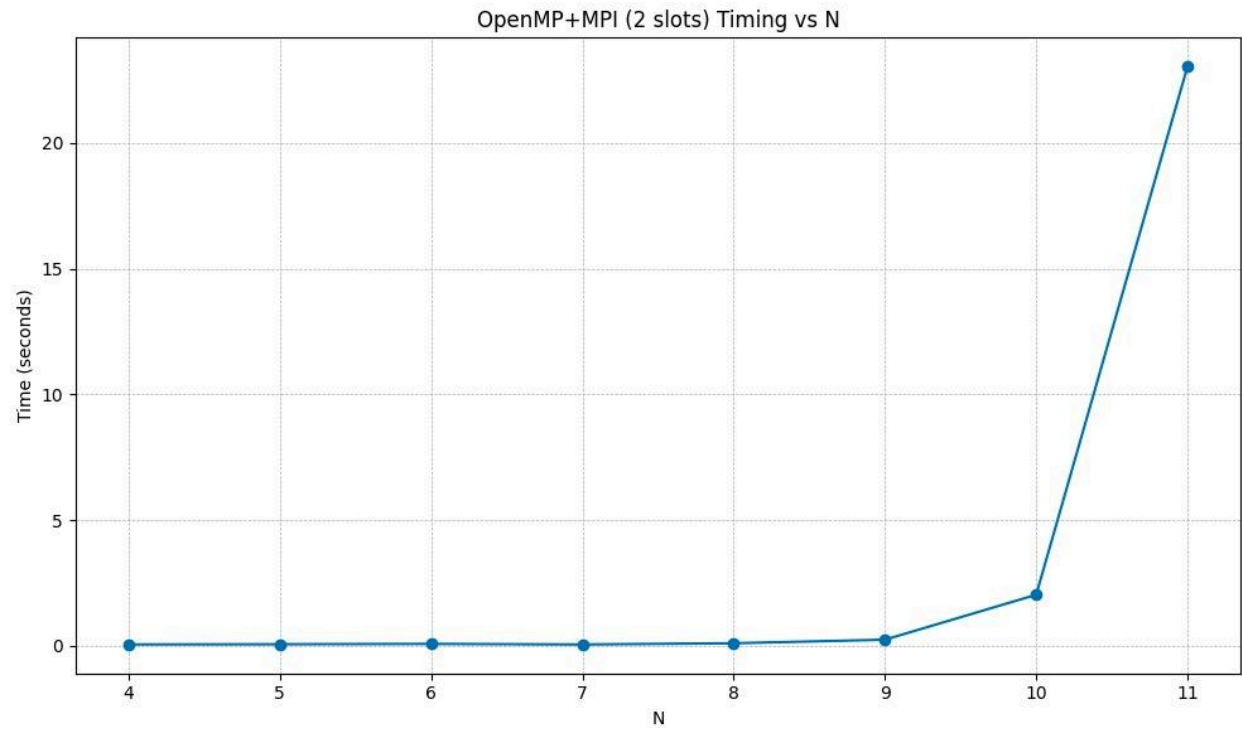
6.3 Chart: Execution Time vs N









There were some storage issues Here.





It can be seen openMP starts doing better from N = 10

6.4 Chart: Speedup Comparison

7. Observations and Challenges

- **High Setup Latency (20s)** observed on the MPI cluster due to networking overhead.
- **Homelab Performance** significantly better, attributed to single-node, high-bandwidth, and large shared memory.
- **I/O Bottlenecks** avoided using single-threaded final write after OpenMP region.

8. Conclusion

This work demonstrates the benefits of hybrid parallelism in computing complex permutation structures. We implemented a memory-optimized, MPI+OpenMP solution to compute parent networks for S_n permutations. The parallel code scales well and enables processing of problem sizes that are infeasible with a sequential approach.

Future work may focus on:

- Reducing I/O bottlenecks using MPI-IO or distributed file systems.
- Implementing a memory-efficient compressed output format.
- Further optimizing inter-process communication using shared memory within nodes.

Appendix

A. Compilation and Execution Instructions

```
mpicxx -O2 -fopenmp parallel_mpi.cpp -o parallel_mpi
```

```
mpirun --hostfile machines -np <num_processes> ./parallel_mpi <N>
```

B. Sample Output File Format

`vertex_id, perm, T1, T2, ..., T(n-1)`

`0, 1-2-3-4, ..., ...`